

# NASM Examples

---

*NASM is a pretty awesome assembler. Let's learn some NASM programming by example. These notes barely scratch the surface of what you can do, so after you've gone through this page, you'll need to hit the [official NASM documentation](#).*

---

## Preliminaries

Please note: all of these examples, except for those in the last few sections, will **only run on a modern 64-bit Linux installation**.

Make sure both nasm and gcc are installed.

## Getting Started

Our first program will use Linux system call 1 to write a message and Linux system call 60 to exit.

hello.asm

```
; -----  
; Writes "Hello, World" to the console using only system calls. Runs on 64-bit Linux only.  
; To assemble and run:  
;  
; nasm -felf64 hello.asm && ld hello.o && ./a.out  
; -----  
  
global _start  
  
section .text  
_start:  
    ; write(1, message, 13)  
    mov     rax, 1          ; system call 1 is write  
    mov     rdi, 1          ; file handle 1 is stdout  
    mov     rsi, message    ; address of string to output  
    mov     rdx, 13         ; number of bytes  
    syscall                ; invoke operating system to do the write  
  
    ; exit(0)  
    mov     eax, 60         ; system call 60 is exit  
    xor     rdi, rdi        ; exit code 0  
    syscall                ; invoke operating system to exit  
message:  
    db      "Hello, World", 10 ; note the newline at the end
```

```
$ nasm -felf64 hello.asm && ld hello.o && ./a.out  
Hello, World
```

## Using a C Library

Remember how in C execution seems to start with "main"? That's because the C library actually has the `_start` label inside itself! The code at `_start` does some initialization, then it calls `main`, then it does some clean up, then it issues system call

60. So you just have to implement `main`. We can do that in assembly:

```

hola.asm
; -----
; Writes "Hola, mundo" to the console using a C library. Runs on Linux or any other system
; that does not use underscores for symbols in its C library. To assemble and run:
;
;   nasm -felf64 hola.asm && gcc hola.o && ./a.out
; -----

    global main
    extern puts

    section .text
main:
    mov     rdi, message        ; This is called by the C library startup code
    call    puts                ; First integer (or pointer) argument in rdi
    ret                                ; puts(message)
                                ; Return from main back into C library wrapper
message:
    db      "Hola, mundo", 0    ; Note strings must be terminated with 0 in C

$ nasm -felf64 hola.asm && gcc hola.o && ./a.out
Hola, mundo

```

## Understanding Calling Conventions

When writing code for 64-bit Linux that integrates with a C library, you must follow the calling conventions, explained fully in the [AMD64 ABI Reference](#). You can also get this information from [Wikipedia](#). The most important points are:

- From left to right, pass as many parameters as will fit in registers. The order in which registers are allocated, are:
  - For integers and pointers, `rdi`, `rsi`, `rdx`, `rcx`, `r8`, `r9`.
  - For floating-point (float, double), `xmm0`, `xmm1`, `xmm2`, `xmm3`, `xmm4`, `xmm5`, `xmm6`, `xmm7`.
- Additional parameters are pushed on the stack, right to left, and are to be *removed by the caller* after the call.
- After the parameters are pushed, the call instruction is made, so when the called function gets control, the return address is at `[rsp]`, the first memory parameter is at `[rsp+8]`, etc.
- **The stack pointer `rsp` must be aligned to a 16-byte boundary before making a call.** Fine, but the process of making a call pushes the return address (8 bytes) on the stack, so when a function gets control, `rsp` is not aligned. You have to make that extra space yourself, by pushing something or subtracting 8 from `rsp`.
- The only registers that the called function is required to preserve (the callee-save registers) are: `rbp`, `rbx`, `r12`, `r13`, `r14`, `r15`. All others are free to be changed by the called function.
- The callee is also supposed to save the control bits of the XCSR and the x87 control word, but x87 instructions are rare in 64-bit code so you probably don't have to worry about this.
- Integers are returned in `rax` or `rdx`: `rax`, and floating point values are returned in `xmm0` or `xmm1`: `xmm0`.

Here is a program that illustrates how registers have to be saved and restored:

```

fib.asm
; -----
; A 64-bit Linux application that writes the first 90 Fibonacci numbers. To
; assemble and run:
;
;   nasm -felf64 fib.asm && gcc fib.o && ./a.out
; -----

    global main
    extern printf

```

```

        section .text
main:
        push    rbx                ; we have to save this since we use it

        mov     ecx, 90            ; ecx will countdown to 0
        xor     rax, rax           ; rax will hold the current number
        xor     rbx, rbx           ; rbx will hold the next number
        inc     rbx                ; rbx is originally 1
print:
        ; We need to call printf, but we are using rax, rbx, and rcx.  printf
        ; may destroy rax and rcx so we will save these before the call and
        ; restore them afterwards.

        push    rax                ; caller-save register
        push    rcx                ; caller-save register

        mov     rdi, format        ; set 1st parameter (format)
        mov     rsi, rax           ; set 2nd parameter (current_number)
        xor     rax, rax           ; because printf is varargs

        ; Stack is already aligned because we pushed three 8 byte registers
        call    printf             ; printf(format, current_number)

        pop     rcx                ; restore caller-save register
        pop     rax                ; restore caller-save register

        mov     rdx, rax           ; save the current number
        mov     rax, rbx           ; next number is now current
        add     rbx, rdx           ; get the new next number
        dec     ecx               ; count down
        jnz     print             ; if not done counting, do some more

        pop     rbx                ; restore rbx before returning
        ret
format:
        db     "%20ld", 10, 0

```

```

$ nasm -felf64 fib.asm && gcc fib.o && ./a.out
0
1
1
2
.
.
.
679891637638612258
1100087778366101931
1779979416004714189

```

## Mixing C and Assembly Language

This program is just a simple function that takes in three integer parameters and returns the maximum value.

maxofthree.asm

```

; -----
; A 64-bit function that returns the maximum value of its three 64-bit integer
; arguments. The function has signature:
;
;   int64_t maxofthree(int64_t x, int64_t y, int64_t z)
;
; Note that the parameters have already been passed in rdi, rsi, and rdx. We
; just have to return the value in rax.
; -----

```

```

global maxofthree
section .text
maxofthree:
    mov     rax, rdi           ; result (rax) initially holds x
    cmp     rax, rsi           ; is x less than y?
    cmovl   rax, rsi           ; if so, set result to y
    cmp     rax, rdx           ; is max(x,y) less than z?
    cmovl   rax, rdx           ; if so, set result to z
    ret                                ; the max will be in rax

```

Here is a C program that calls the assembly language function.

#### callmaxofthree.c

```

/*
 * A small program that illustrates how to call the maxofthree function we wrote in
 * assembly language.
 */

#include <stdio.h>
#include <inttypes.h>

int64_t maxofthree(int64_t, int64_t, int64_t);

int main() {
    printf("%ld\n", maxofthree(1, -4, -7));
    printf("%ld\n", maxofthree(2, -6, 1));
    printf("%ld\n", maxofthree(2, 3, 1));
    printf("%ld\n", maxofthree(-2, 4, 3));
    printf("%ld\n", maxofthree(2, -6, 5));
    printf("%ld\n", maxofthree(2, 4, 6));
    return 0;
}

```

```

$ nasm -felf64 maxofthree.asm && gcc callmaxofthree.c maxofthree.o && ./a.out
1
2
3
4
5
6

```

## Command Line Arguments

You know that in C, `main` is just a plain old function, and it has a couple parameters of its own:

```
int main(int argc, char** argv)
```

Here is a program that uses this fact to simply echo the commandline arguments to a program, one per line:

#### echo.asm

```

; -----
; A 64-bit program that displays its command line arguments, one per line.
;
; On entry, rdi will contain argc and rsi will contain argv.
; -----

global main
extern puts
section .text
main:
    push    rdi                ; save registers that puts uses
    push    rsi
    sub     rsp, 8              ; must align stack before call

```

```

    mov     rdi, [rsi]           ; the argument string to display
    call    puts                ; print it

    add     rsp, 8               ; restore %rsp to pre-aligned value
    pop     rsi                 ; restore registers puts used
    pop     rdi

    add     rsi, 8               ; point to next argument
    dec     rdi                 ; count down
    jnz     main                ; if not done counting keep going

    ret

```

```

$ nasm -felf64 echo.asm && gcc echo.o && ./a.out dog 22 -zzz "hi there"
./a.out
dog
22
-zzz
hi there

```

## A Longer Example

Note that as far as the C Library is concerned, command line arguments are always strings. If you want to treat them as integers, call `atoi`. Here's a neat program to compute  $x^y$ .

power.asm

```

; -----
; A 64-bit command line application to compute x^y.
;
; Syntax: power x y
; x and y are (32-bit) integers
; -----

    global  main
    extern  printf
    extern  puts
    extern  atoi

    section .text
main:
    push    r12                ; save callee-save registers
    push    r13
    push    r14
    ; By pushing 3 registers our stack is already aligned for calls

    cmp     rdi, 3              ; must have exactly two arguments
    jne     error1

    mov     r12, rsi            ; argv

    ; We will use ecx to count down from the exponent to zero, esi to hold the
    ; value of the base, and eax to hold the running product.

    mov     rdi, [r12+16]       ; argv[2]
    call    atoi                ; y in eax
    cmp     eax, 0              ; disallow negative exponents
    jl      error2
    mov     r13d, eax            ; y in r13d

    mov     rdi, [r12+8]        ; argv
    call    atoi                ; x in eax
    mov     r14d, eax            ; x in r14d

```

```

    mov     eax, 1                ; start with answer = 1
check:
    test    r13d, r13d           ; we're counting y downto 0
    jz      gotit                ; done
    imul    eax, r14d            ; multiply in another x
    dec     r13d
    jmp     check
gotit:
                                ; print report on success
    mov     rdi, answer
    movsxd  rsi, eax
    xor     rax, rax
    call    printf
    jmp     done
error1:
                                ; print error message
    mov     edi, badArgumentCount
    call    puts
    jmp     done
error2:
                                ; print error message
    mov     edi, negativeExponent
    call    puts
done:
                                ; restore saved registers
    pop     r14
    pop     r13
    pop     r12
    ret

answer:
    db      "%d", 10, 0
badArgumentCount:
    db      "Requires exactly two arguments", 10, 0
negativeExponent:
    db      "The exponent may not be negative", 10, 0

```

```

$ nasm -felf64 power.asm && gcc -o power power.o
$ ./power 2 19
524288
$ ./power 3 -8
The exponent may not be negative
$ ./power 1 500
1
$ ./power 1
Requires exactly two arguments

```

## Floating Point Instructions

Floating-point arguments go into the xmm registers. Here is a simple function for summing the values in a double array:

sum.asm

```

; -----
; A 64-bit function that returns the sum of the elements in a floating-point
; array. The function has prototype:
;
; double sum(double[] array, uint64_t length)
; -----

    global sum
    section .text

sum:
    xorsd   xmm0, xmm0          ; initialize the sum to 0
    cmp     rsi, 0              ; special case for length = 0
    je      done

next:
    addsd   xmm0, [rdi]         ; add in the current array element
    add     rdi, 8              ; move to next array element
    dec     rsi                ; count down

```

```

        jnz     next                ; if not done counting, continue
done:
        ret                        ; return value already in xmm0

```

A C program that calls it:

#### callsum.c

```

/*
 * Illustrates how to call the sum function we wrote in assembly language.
 */

#include <stdio.h>
#include <inttypes.h>

double sum(double[], uint64_t);

int main() {
    double test[] = {
        40.5, 26.7, 21.9, 1.5, -40.5, -23.4
    };
    printf("%20.7f\n", sum(test, 6));
    printf("%20.7f\n", sum(test, 2));
    printf("%20.7f\n", sum(test, 0));
    printf("%20.7f\n", sum(test, 3));
    return 0;
}

```

```

$ nasm -felf64 sum.asm && gcc sum.o callsum.c && ./a.out
      26.7000000
      67.2000000
       0.0000000
      89.1000000

```

## Data Sections

The text section is read-only on most operating systems, so you might find the need for a data section. On most operating systems, the data section is only for initialized data, and you have a special .bss section for uninitialized data. Here is a program that averages the command line arguments, expected to be integers, and displays the result as a floating point number.

#### average.asm

```

; -----
; 64-bit program that treats all its command line arguments as integers and
; displays their average as a floating point number. This program uses a data
; section to store intermediate results, not that it has to, but only to
; illustrate how data sections are used.
; -----

        global  main
        extern  atoi
        extern  printf
        default rel

        section .text
main:
        dec     rdi                ; argc-1, since we don't count program name
        jz      nothingToAverage
        mov     [count], rdi       ; save number of real arguments

accumulate:
        push    rdi                ; save register across call to atoi
        push    rsi
        mov     rdi, [rsi+rdi*8]   ; argv[rdi]
        call    atoi              ; now rax has the int value of arg

```

```

    pop    rsi                ; restore registers after atoi call
    pop    rdi
    add     [sum], rax         ; accumulate sum as we go
    dec     rdi                ; count down
    jnz     accumulate        ; more arguments?
average:
    cvtsi2sd xmm0, [sum]
    cvtsi2sd xmm1, [count]
    divsd   xmm0, xmm1        ; xmm0 is sum/count
    mov     rdi, format        ; 1st arg to printf
    mov     rax, 1             ; printf is varargs, there is 1 non-int argument

    sub     rsp, 8             ; align stack pointer
    call    printf             ; printf(format, sum/count)
    add     rsp, 8             ; restore stack pointer

    ret

nothingToAverage:
    mov     rdi, error
    xor     rax, rax
    call    printf
    ret

section .data
count: dq 0
sum: dq 0
format: db "%g", 10, 0
error: db "There are no command line arguments to average", 10, 0

```

```

$ nasm -felf64 average.asm && gcc average.o && ./a.out 19 8 21 -33
3.75
$ nasm -felf64 average.asm && gcc average.o && ./a.out
There are no command line arguments to average

```

## Recursion

Perhaps surprisingly, there's nothing out of the ordinary required to implement recursive functions. You just have to be careful to save registers, as usual.

### factorial.asm

```

; -----
; An implementation of the recursive function:
;
;  uint64_t factorial(uint64_t n) {
;      return (n <= 1) ? 1 : n * factorial(n-1);
;  }
; -----

global factorial

section .text
factorial:
    cmp     rdi, 1             ; n <= 1?
    jnbe    L1                 ; if not, go do a recursive call
    mov     rax, 1             ; otherwise return 1
    ret

L1:
    push    rdi                ; save n on stack (also aligns %rsp!)
    dec     rdi                ; n-1
    call    factorial           ; factorial(n-1), result goes in %rax
    pop     rdi                ; restore n
    imul    rax, rdi           ; n * factorial(n-1), stored in %rax
    ret

```



An example caller:

#### callfactorial.c

```
/*
 * An application that illustrates calling the factorial function defined elsewhere.
 */

#include <stdio.h>
#include <inttypes.h>

uint64_t factorial(uint64_t n);

int main() {
    for (uint64_t i = 0; i < 20; i++) {
        printf("factorial(%2lu) = %lu\n", i, factorial(i));
    }
    return 0;
}
```

```
$ nasm -felf64 factorial.asm && gcc -std=c99 factorial.o callfactorial.c && ./a.out
factorial( 0) = 1
factorial( 1) = 1
factorial( 2) = 2
factorial( 3) = 6
factorial( 4) = 24
factorial( 5) = 120
factorial( 6) = 720
factorial( 7) = 5040
factorial( 8) = 40320
factorial( 9) = 362880
factorial(10) = 3628800
factorial(11) = 39916800
factorial(12) = 479001600
factorial(13) = 6227020800
factorial(14) = 87178291200
factorial(15) = 1307674368000
factorial(16) = 20922789888000
factorial(17) = 355687428096000
factorial(18) = 6402373705728000
factorial(19) = 121645100408832000
```

## SIMD Parallelism

The XMM registers can do arithmetic on floating point values one operation at a time or multiple operations at a time. The operations have the form:

operation xmmregister\_or\_memorylocation, xmmregister

For floating point addition, the instructions are:

addpd – do 2 double-precision additions  
 addps – do just one double-precision addition, using the low 64-bits of the register  
 addsd – do 4 single-precision additions  
 addss – do just one single-precision addition, using the low 32-bits of the register

TODO - show a function that processes an array of floats, 4 at a time.

## Saturated Arithmetic

The XMM registers can also do arithmetic on integers. The instructions have the form:

```
operation    xmmregister_or_memorylocation, xmmregister
```

For integer addition, the instructions are:

```
paddb  — do 16 byte additions
paddw  — do 8 word additions
paddd  — do 4 dword additions
paddq  — do 2 qword additions
paddsb — do 16 byte additions with signed saturation (80..7F)
paddsw — do 8 word additions with unsigned saturation (8000..7FFF)
paddusb — do 16 byte additions with unsigned saturation (00..FF)
paddusw — do 8 word additions with unsigned saturation (00..FFFF)
```

TODO - SHOW AN EXAMPLE

## Graphics

TODO

## Local Variables and Stack Frames

First, please read [Eli Bendersky's article](#) That overview is more complete than my brief notes.

When a function is called the caller will first put the parameters in the correct registers then issue the `call` instruction. Additional parameters beyond those covered by the registers will be pushed on the stack prior to the call. The call instruction puts the return address on the top of stack. So if you have the function

```
int64_t example(int64_t x, int64_t y) {
    int64_t a, b, c;
    b = 7;
    return x * b + y;
}
```

Then on entry to the function, `x` will be in `edi`, `y` will be in `esi`, and the return address will be on the top of the stack. Where can we put the local variables? An easy choice is on the stack itself, though if you have enough registers, use those.

If you are running on a machine that respect the standard ABI, you can leave `rsp` where it is and access the "extra parameters" and the local variables directly from `rsp` for example:

```

+-----+
rsp-24 |    a    |
+-----+
rsp-16 |    b    |
+-----+
rsp-8  |    c    |
+-----+
rsp   | retaddr |
+-----+
rsp+8 | caller's |
      | stack   |
      | frame   |
      | ...     |
+-----+
```

So our function looks like this:

```

global example
section .text
example:
    mov     qword [rsp-16], 7
    mov     rax, rdi
    imul    rax, [rsp+8]
    add     rax, rsi
    ret

```

If our function were to make another call, you would have to adjust `rsp` to get out of the way at that time.

On Windows you can't use this scheme because if an interrupt were to occur, everything above the stack pointer gets plastered. This doesn't happen on most other operating systems because there is a "red zone" of 128 bytes past the stack pointer which is safe from these things. In this case, you can make room on the stack immediately:

```

example:
    sub     rsp, 24

```

so our stack looks like this:

```

      +-----+
rsp   |      a      |
      +-----+
rsp+8 |      b      |
      +-----+
rsp+16|      c      |
      +-----+
rsp+24| retaddr    |
      +-----+
rsp+32| caller's   |
      | stack      |
      | frame      |
      | ...        |
      +-----+

```

Here's the function now. Note that we have to remember to replace the stack pointer before returning!

```

global example
section .text
example:
    sub     rsp, 24
    mov     qword [rsp+8], 7
    mov     rax, rdi
    imul    rax, [rsp+8]
    add     rax, rsi
    add     rsp, 24
    ret

```

## Using NASM on OS X

TODO

## Using NASM on Windows

TODO